WorldForge 2 - Open Issues and Release-First Improvement Plan

Written: 2026-10-15
Status: ACTIVE. This is the current top priority for all work in this repo.
Source: a /court trial of four claims (12 independent juror agents per claim, prosecution, defense, judge), each finding then checked against the repository.
Verdicts: docs/court/2026-10-15/
PDF copy: docs/IMPROVEMENT_PLAN.pdf (regenerate with: python scripts/plan_to_pdf.py)

1. The verdicts

Score = jurors who voted FOR the claim, out of 12, scaled to 10.

- Claim 1 - "Worth continuing in its current direction": GUILTY 11-1. Score 1/10.
- Claim 2 - "Ready (or close) to sell": GUILTY 12-0. Score 0/10.
- Claim 3 - "Generator + deep campaign in one app is the right scope": GUILTY 9-3. Score 2.5/10.
- Claim 4 - "Tauri + Rust tiles + SQLite is the right foundation": NOT GUILTY 9-3. Score 7.5/10.

Bottom line, the same in all four rulings: the engine is sound, the ORDER OF WORK is wrong. Nobody argued for a rewrite or a new stack. The project keeps adding simulation layers while nothing has shipped, no outsider has used it, and the parts a buyer would touch first (install, saved files, memory on their machine) are unproven.

Redesign question: the court found no case for redesigning the program. The only "new idea" it supports is a reframing of the PRODUCT, not the code: sell it as a science-grounded worldbuilding map tool, with the living campaign as a second mode, and let testers decide how deep the campaign goes.

2. Open issues

I-1 No evidence anyone wants it [severity: critical]

- No users, testers, buyers, price, sales page or waitlist anywhere.
- All 12 jurors on claim 2 rested their vote on this.

I-2 Nothing can be installed or bought [critical]

- No LICENSE file, no release tags, version still 0.1.0 in package.json and Cargo.toml.
- README only documents dev mode (npm run tauri dev, run.bat).
- CI (.github/workflows/ci.yml) type-checks, tests and runs vite build, but never produces a Tauri installer.

I-3 Features are being added before anything ships [critical]

- Sept-Oct 2026 went into Living World rows 02-10 (face genomes, councils, leisure and games, masterworks) and a UI redesign.
- Each new layer adds code for one person to maintain and new economy side effects (see I-9).

I-4 Saved worlds and campaigns [high]

- FIXED for new saves: slice 1 of docs/TWO_APPS_AND_FILE_UPLOAD_PLAN.md (section 8) stopped save_world_as and new_campaign from deleting the world's human layer. Gate: campaign_key_split_tests.
- STILL OPEN: .worldforge files written BEFORE that fix cannot start a campaign without their matching .campaign file (the recovery card in ChroniclePanel explains this).
- STILL OPEN: slices 3-5 (.wf2 bundle, sibling-file pairing, drag-and-drop) are not built.
- STILL OPEN: no test opens a saved file from an OLDER build and begins a campaign. SCOREBOARD.md itself says "a compatibility claim with no old-save fixture behind it is a hope".

I-5 Peak memory has never been measured [high]

- SCOREBOARD.md: "Time is benchmarked; memory is not, and memory is the likelier failure on a customer's machine."
- The year-30 out-of-memory crash (unbounded history journal) is recorded as fixed, but no long full-size campaign has been measured since.
- All 9 jurors who acquitted the architecture said they would flip if memory turns out bad.

I-6 Campaign speed at real scale [medium]

- CLAUDE.md section 5.5: at ~1,200 hubs the trade phase is still ~8,300-10,200 ms per simulated year after the house_for cache (down from 18-26 s).
- A decade-long campaign on a big world is minutes of waiting. Fine for testers, worth fixing before a paid release.

I-7 Product identity is undecided [high]

- Still open: one app, or a world generator plus a campaign player?
- There is no one-sentence pitch. The Designer juror: "a stranger could not say in five seconds whether this is a map generator or a merchant-and-dynasty sim".

I-8 Legal basics for a paid product [high before any sale]

- No commercial licence or end-user agreement (EULA).
- No audit of dependency licences (Tauri, React, PixiJS, rusqlite and the rest of the npm/cargo tree).
- No check of data rights: the Koppen-Geiger reference map used by the climate harness, the commodity-history text in the Goods Codex, any bundled fonts or images.
- Most commits (728 of 988) were written by an AI agent. Check how your jurisdiction treats copyright in AI-assisted code before selling.
- This list is not legal advice. Get a short professional review before the first sale.

I-9 Quality debt [medium]

- The frontend (~33k lines per SCOREBOARD, ~60k TS/TSX in total) has zero tests.
- Rust <-> TypeScript types are mirrored by hand (types/campaign.ts), so a renamed field fails silently.
- The top-10% wealth share has fallen out of its band, cause "not yet explained" (SCOREBOARD 2026-09-30b). The bisect is queued.
- Tests have been red on main at times (inheritance gate, unrest_topples_councils). Check their current state in CI.
- There are 116 git branches. Prune the merged ones.

I-10 Local test environment [low]

- cargo test binaries fail to launch on the owner's Windows machine (STATUS_ENTRYPOINT_NOT_FOUND). Use cargo check locally and rely on CI for tests.

3. The plan

Work the phases in order. Each item has a "done when" so progress is checkable. Tick items here ([x]) as they land, and regenerate the PDF.

Phase 0 - Feature freeze (now)

- **DONE** Freeze written into CLAUDE.md section 0: no new Living World / people-layer / campaign features until v0.1 has shipped to testers. Bug fixes, performance, save safety and release work are allowed.
- **TODO** Decide with the owner what happens to unfinished Living World rows: park them on a branch or leave them as-is. Do not extend them.

Phase 1 - Save safety (I-4)

- **TODO** Commit an old-save fixture: a small .worldforge (+ .campaign) written by an older build, under src-tauri/tests or test fixtures.
- **TODO** Regression test: open the fixture, then begin a campaign, and assert the campaign has hubs and runs one tick. Done when it is green in CI.
- **TODO** Slice 4 - sibling-file pairing: opening X.worldforge offers X.campaign next to it automatically.
- **TODO** Optional, after testers: slice 3 (.wf2 bundle) and slice 5 (drag-and-drop).

Phase 2 - Measure (I-5, I-6)

- **TODO** Before measuring, write down the target machine (for example 8 GB RAM laptop) and a peak-memory limit.
- **TODO** Profile one campaign on a 3600x1800 world, well past year 30 (aim for year 60), recording peak memory (working set) and seconds per simulated year. WF2_PROFILE=1 already gives the per-phase time.
- **TODO** Record the numbers as a new row in docs/SCOREBOARD.md.
- **TODO** If memory is over the limit, find the growing structure (same method as the journal fix). If over budget because of the tile/WorldBuffer design itself, put claim 4 back on trial.
- **TODO** Next target for speed: shipment logic in dispatch() (CLAUDE.md section 5.5 names it as the largest remaining cost).

Phase 3 - Decide the product (I-7)

- **TODO** One app or two. The court's lean: ONE app, map tool first, campaign as a mode. Record the decision at the top of docs/TWO_APPS_AND_FILE_UPLOAD_PLAN.md.
- **TODO** Write the one-sentence pitch and put it as the first line of README.md.
- **TODO** Name the first buyer (for example: fantasy authors, TTRPG game masters, worldbuilding hobbyists).

Phase 4 - Legal basics (I-8)

- **TODO** Choose a licence model (closed commercial, or source-available) and add LICENSE.
- **TODO** Generate a dependency licence report (cargo-about or cargo-license for Rust, license-checker for npm) and save it in docs/.
- **TODO** Check the licence of the Koppen-Geiger dataset and the commodity-history text, and replace anything that cannot be sold.
- **TODO** Draft a short EULA. Have it, and the AI-code question, reviewed once.

Phase 5 - Ship v0.1 to testers (I-1, I-2)

- **TODO** Add a CI job that runs tauri build on Windows and uploads the installer as an artifact.
- **TODO** Bump to 0.1.0-beta, tag v0.1.0-beta, attach the installer to a GitHub release (private or public).
- **TODO** Add an "Install" section to the README that does not need Node or Rust.
- **TODO** Recruit 5-10 outside worldbuilders. Ask each to: install, make a world, export a map, save, reopen, start a campaign, run 10 years.
- **TODO** Collect what broke, what they used, what they ignored, and whether they would pay (and how much).
- **TODO** Put up a one-page waitlist or pre-order page with a stated price, and count sign-ups.

Phase 6 - Let the testers steer

- **TODO** Only now decide whether the campaign / people layers grow, stay, or get cut, based on what testers actually used.

- **TODO** Re-run /court on claims 1-3 with the tester evidence in the case file.

Housekeeping (any time, small)

- **TODO** Fix or explain the red tests on main, bisect the top-10% wealth share drift.
- **TODO** Prune merged branches (116 now).
- **TODO** Add a generated Rust -> TS type check (ts-rs or specta) for the campaign types, or at least one smoke test per panel.

4. What NOT to do

- Do not rewrite the stack. Claim 4 was acquitted: every recorded failure was in application code, not in Tauri, Rust, tiles or SQLite.
- Do not split into two codebases before testers have used it.
- Do not start new Living World rows, new goods, new institutions or new panels until Phase 5 is done.
- Do not treat "the tests pass" as "it is ready". The tests measure the code, not demand.

5. How to use and update this file

- Every new Claude Code session: read CLAUDE.md section 0 first, then this file. Tell the owner which phase is current and which items are open before starting other work.
- If the owner asks for a new feature while Phase 0-5 items are open, remind them of the freeze and of the next open item, then do what they decide.
- When an item lands: tick it here, add evidence (commit, test name, number), and regenerate the PDF: `python scripts/plan_to_pdf.py`
- When all of Phase 5 is ticked, the freeze is over. Update CLAUDE.md section 0.